

## Movie\_recommendation

Cmd 1

```
1 print "sc type: ", type(sc)
2 print "Driver Program name: ", sc.appName
3 print "Spark version: ", sc.version
```

sc type: <class '\_\_main\_\_.RemoteContext'>

Driver Program name: Databricks Shell

Spark version: 2.1.1

Command took 0.16 seconds -- by meghana.rwgsq@gmail.com at 6/10/2017, 2:24:41 AM on

Movies\_Project  
Cmd 2

```
1 moviesDF = spark.sql("select * from movies")
2 moviesDF.show(truncate = False)
```

### ► (1) Spark Jobs

```
+-----+-----+-----+-----+
+-----+
|movieId|title                                     |genres
|
+-----+-----+-----+-----+
+-----+
|1       |Toy Story (1995)                         |Adventure|Animation|Children|Come
dy|Fantasy|
|2       |Jumanji (1995)                           |Adventure|Children|Fantasy
|
|3       |Grumpier Old Men (1995)                  |Comedy|Romance
|
|4       |Waiting to Exhale (1995)                 |Comedy|Drama|Romance
|
|5       |Father of the Bride Part II (1995)      |Comedy
|
|6       |Heat (1995)                             |Action|Crime|Thriller
|
|7       |Sabrina (1995)                          |Comedy|Romance
|
|8       |Tom and Huck (1995)                     |Adventure|Children
|
|9       |Sudden Death (1995)                     |Action
|
|10      |GoldenEye (1995)                        |Action|Adventure|Thriller
|
|11      |American President, The (1995)          |Comedy|Drama|Romance
|
|12      |Dracula: Dead and Loving It (1995)      |Comedy|Horror
```

13	Balto (1995)	Adventure Animation Children
14	Nixon (1995)	Drama
15	Cutthroat Island (1995)	Action Adventure Romance
16	Casino (1995)	Crime Drama
17	Sense and Sensibility (1995)	Drama Romance
18	Four Rooms (1995)	Comedy
19	Ace Ventura: When Nature Calls (1995)	Comedy
20	Money Train (1995)	Action Comedy Crime Drama Thriller

only showing top 20 rows

Command took 3.27 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 10:14:36 AM on

movies\_project  
Cmd 3

```
1 spark.sql("SELECT count(*) FROM movies").show()
```

#### ► (1) Spark Jobs

```
+-----+
|count(1)|
+-----+
|   27278|
+-----+
```

Command took 1.31 seconds -- by meghana.rwgsql@gmail.com at 8/30/2017, 9:09:07 PM on

movies\_project  
Cmd 4

```
1 moviesDF.printSchema()
```

```
root
|-- movieId: string (nullable = true)
|-- title: string (nullable = true)
|-- genres: string (nullable = true)
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:50:08 PM on

movies\_project  
Cmd 5

```
1 ratingsDF = spark.sql("select * from ratings")
2 ratingsDF.show()
```

► (1) Spark Jobs

```
+-----+-----+-----+-----+
|userId|movieId|rating| timestamp|
+-----+-----+-----+-----+
|      1|      2|    3.5|1112486027|
|      1|     29|    3.5|1112484676|
|      1|     32|    3.5|1112484819|
|      1|     47|    3.5|1112484727|
|      1|     50|    3.5|1112484580|
|      1|    112|    3.5|1094785740|
|      1|    151|    4.0|1094785734|
|      1|    223|    4.0|1112485573|
|      1|    253|    4.0|1112484940|
|      1|    260|    4.0|1112484826|
|      1|    293|    4.0|1112484703|
|      1|    296|    4.0|1112484767|
|      1|    318|    4.0|1112484798|
|      1|    337|    3.5|1094785709|
|      1|    367|    3.5|1112485980|
|      1|    541|    4.0|1112484603|
|      1|    589|    3.5|1112485557|
|      1|    593|    3.5|1112484661|
|      1|    653|    3.0|1094785691|
|      1|    919|    3.5|1094785621|
+-----+-----+-----+-----+
```

only showing top 20 rows

Command took 0.43 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:28 AM on  
 movies\_project  
 Cmd 6

```
1 spark.sql("SELECT count(*) FROM ratings").show()
```

► (1) Spark Jobs

```
+-----+
|count(1)|
+-----+
|20000263|
+-----+
```

Command took 8.88 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 9:09:32 PM on  
 movies\_project  
 Cmd 7

```
1 ratingsDF.printSchema()
```

```
root
```

```
|-- userId: string (nullable = true)
|-- movieId: string (nullable = true)
|-- rating: string (nullable = true)
|-- timestamp: string (nullable = true)
```

Command took 0.07 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 12:25:52 AM on

Movies\_Project  
Cmd 8

```
1 ratingsDF = ratingsDF.withColumn('rating',
  ratingsDF["rating"].cast("float"))
2 ratingsDF = ratingsDF.withColumn('userId', ratingsDF["userId"].cast("int"))
3 ratingsDF = ratingsDF.withColumn('movieId',
  ratingsDF["movieId"].cast("int"))
4 ratingsDF.show()
```

#### ► (1) Spark Jobs

```
+-----+-----+-----+-----+
|userId|movieId|rating| timestamp|
+-----+-----+-----+-----+
|    1|      2|    3.5|1112486027|
|    1|     29|    3.5|1112484676|
|    1|     32|    3.5|1112484819|
|    1|     47|    3.5|1112484727|
|    1|     50|    3.5|1112484580|
|    1|    112|    3.5|1094785740|
|    1|    151|    4.0|1094785734|
|    1|    223|    4.0|1112485573|
|    1|    253|    4.0|1112484940|
|    1|    260|    4.0|1112484826|
|    1|    293|    4.0|1112484703|
|    1|    296|    4.0|1112484767|
|    1|    318|    4.0|1112484798|
|    1|    337|    3.5|1094785709|
|    1|    367|    3.5|1112485980|
|    1|    541|    4.0|1112484603|
|    1|    589|    3.5|1112485557|
|    1|    593|    3.5|1112484661|
|    1|    653|    3.0|1094785691|
|    1|    919|    3.5|1094785621|
+-----+-----+-----+-----+
```

only showing top 20 rows

Command took 0.53 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:37 AM on

movies\_project  
Cmd 9

```
1 from pyspark.sql.functions import from_unixtime
2 ratingsDF = ratingsDF.withColumn('timestamp',
  from_unixtime(ratingsDF.timestamp))
3
```

Command took 0.12 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:42 AM on

movies\_project  
Cmd 10

```
1 from pyspark.sql.types import TimestampType
2 ratingsDF = ratingsDF.withColumn('timestamp',
  ratingsDF.timestamp.cast(TimestampType()))
```

Command took 0.12 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:46 AM on

movies\_project  
Cmd 11

```
1 ratingsDF.printSchema()
```

root

```
-- userId: integer (nullable = true)
-- movieId: integer (nullable = true)
-- rating: float (nullable = true)
-- timestamp: timestamp (nullable = true)
```

Command took 0.07 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:51 AM on

movies\_project  
Cmd 12

```
1 moviesDF.describe().show()
```

#### ► (1) Spark Jobs

summary	movieId	title	genres
count	27278	27278	27278
mean	59855.48057042305	null	null
stddev	44429.31469707313	null	null
min	1	""Great Performa...	(no genres listed)
max	99999	貞子3D (2012)	Western

Command took 2.32 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:17:55 AM on

movies\_project  
Cmd 13

```

1 #m_DF =
  sqlContext.read.format("csv").load("/FileStore/tables/4riqczku1497073456815
  /movies.csv")
2 #r_df =
  sqlContext.read.format("csv").load("/FileStore/tables/f9oap0qt1497073661552
  /ratings.csv")

```

## ► (2) Spark Jobs

Command took 0.93 seconds -- by meghana.rwgsql@gmail.com at 6/10/2017, 2:37:14 AM on

Movies\_Project  
Cmd 14

```
1 ratingsDF.describe().show()
```

## ► (1) Spark Jobs

summary	userId	movieId	rating
count	20000263	20000263	20000263
mean	69045.87258292554	9041.567330339605	3.5255285642993797
stddev	40038.62665316182	19789.47744541297	1.05198891929425
min	1	1	0.5
max	138493	131262	5.0

Command took 14.05 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 10:18:02 AM on

Movies\_project  
Cmd 15

```

1 #To check if Dataframe contains None/blank values
2 print(ratingsDF[ratingsDF.rating == None].count())
3 print(ratingsDF[ratingsDF.userId == None].count())
4 print(ratingsDF[ratingsDF.movieId == None].count())

```

## ► (3) Spark Jobs

0  
0  
0

Command took 23.31 seconds -- by meghana.rwgsql@gmail.com at 7/20/2017, 5:23:05 PM on

Movies\_Project  
Cmd 16

```

1 #Counting number of users who gave highest ratings to the movie
2 num_users = spark.sql("SELECT count(userId) as no_of_user, ratings.movieId,
   movies.title, rating FROM ratings JOIN movies on movies.movieId =
   ratings.movieId Group by ratings.movieId, movies.title, rating Order by
   rating DESC LIMIT 10")
3 num_users.show(truncate = False)
4

```

## ► (2) Spark Jobs

no_of_user	movieId	title	rating
83	1005	D3: The Mighty Ducks (1996)	5.0
3	7878	Straight to Hell (1987)	5.0
91	688	Operation Dumbo Drop (1995)	5.0
2536	1517	Austin Powers: International Man of Mystery (1997)	5.0
102	6059	Recruit, The (2003)	5.0
231	88810	Help, The (2011)	5.0
210	1888	Hope Floats (1998)	5.0
8272	924	2001: A Space Odyssey (1968)	5.0
13	6454	Music Box (1989)	5.0
94	118	If Lucy Fell (1996)	5.0

Command took 19.46 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:26:26 AM on

movies\_project  
Cmd 17

```

1 #Retrieving the highest average ratings for each movies
2 Movies_avg_ratings = spark.sql("SELECT ratings.movieId, movies.title,
   avg(rating) as avg_rating FROM ratings JOIN movies on movies.movieId =
   ratings.movieId \
3                                     Group by ratings.movieId, movies.title Order by
   avg(rating) DESC LIMIT 20")
4 Movies_avg_ratings.show(truncate = False)

```

## ► (2) Spark Jobs

movieId	title	avg_rating
113790	Peace, Propaganda & the Promised Land (2004)	5.0
111548	Welcome to Australia (1999)	5.0
129295	A Gun for Jennifer (1997)	5.0
116387	Muddy River (1981)	5.0
94949	Boy Meets Boy (2008)	5.0
94431	Ella Lola, a la Trilby (1898)	5.0

130996	The Beautiful Story (1992)	5.0	
109253	Argentina latente (2007)	5.0	
32230	Snow Queen, The (Lumikuningatar) (1986)	5.0	
103753	Human Behavior Experiments, The (2006)	5.0	
89133	Boys (Drenge) (1977)	5.0	
100830	Blue Swallow (Cheong yeon) (2005)	5.0	
106082	Shock and Awe: The Story of Electricity (2011)	5.0	
26718	Life On A String (Bian chang Bian Zou) (1991)	5.0	
94737	Boys Diving, Honolulu (1901)	5.0	
131050	Stargate SG-1 Children of the Gods - Final Cut (2009)	5.0	
103143	Donos de Portugal (2012)	5.0	
121029	No Distance Left to Run (2010)	5.0	
117061	The Green (2011)	5.0	
108527	Catastroika (2012)	5.0	
+-----+-----+-----+-----+			

Command took 21.88 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:26:53 AM on

movies\_project  
Cmd 18

```

1 #Counts number of movies rated by each user
2 num_Movies_rated_by_user = spark.sql("SELECT count(movieId) as
   no_of_movies, userId from ratings GROUP BY userId order by userId")
3 num_Movies_rated_by_user.show()
4
5
```

#### ► (1) Spark Jobs

+-----+-----+		
no_of_movies	userId	
+-----+-----+		
	175	1
	38	10
	52	100
	57	1000
	73	10000
	103	100000
	21	100001
	56	100002
	99	100003
	186	100004
	38	100005
	116	100006
	90	100007
	141	100008
	455	100009
	709	10001



```

|          35|100010|
|          152|100011|
|           20|100012|
|          226|100013|
+-----+-----+
only showing top 20 rows

```

Command took 11.69 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 4:15:03 PM on

Movies\_Recommendation  
Cmd 19

```

1 #Counting distinct users and movies from ratings table
2 numUsers = spark.sql("SELECT count(Distinct userId) as no_of_Users from
  ratings")
3 numMovies = spark.sql("SELECT count(Distinct movieId) as no_of_Movies from
  ratings")
4 numUsers.show()
5 numMovies.show()

```

## ► (2) Spark Jobs

```

+-----+
|no_of_Users|
+-----+
|      138493|
+-----+

+-----+
|no_of_Movies|
+-----+
|       26744|
+-----+

```

Command took 24.14 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:27:30 AM on

movies\_project  
Cmd 20

```

1 ratings_count = ratingsDF.count()
2 movies_count = moviesDF.count()
3
4 print 'There are %s ratings and %s movies in the datasets' %
  (ratings_count, movies_count)
5 print 'Ratings:'
6 ratingsDF.show(3)
7 print 'Movies:'
8 moviesDF.show(3, truncate=False)
9
10 assert ratings_count == ratings_count
11 assert movies_count == movies_count

```

#### ► (4) Spark Jobs

There are 20000263 ratings and 27278 movies in the datasets

Ratings:

```
+-----+-----+-----+-----+
|userId|movieId|rating|          timestamp|
+-----+-----+-----+-----+
|      1|      2|    3.5|2005-04-02 23:53:47|
|      1|     29|    3.5|2005-04-02 23:31:16|
|      1|     32|    3.5|2005-04-02 23:33:39|
+-----+-----+-----+-----+
```

only showing top 3 rows

Movies:

```
+-----+-----+-----+-----+-----+-----+
|movieId|title                |genres
+-----+-----+-----+-----+-----+-----+
|1       |Toy Story (1995)      |Adventure|Animation|Children|Comedy|Fantasy|
|2       |Jumanji (1995)        |Adventure|Children|Fantasy
|3       |Grumpier Old Men (1995)|Comedy|Romance
+-----+-----+-----+-----+-----+-----+
```

only showing top 3 rows

Command took 8.93 seconds -- by meghana.rwgsq1@gmail.com at 8/31/2017, 10:30:18 AM on

movies\_project  
Cmd 21

```

1 #MOVIES WITH HIGHEST AVERAGE RATINGS
2
3 from pyspark.sql import functions as F
4
5 # Rename column from movieId to ID
6 movies_DF = moviesDF.withColumnRenamed('movieId', 'ID')
7
8 # From ratingsDF, create a movie_ids_with_avg_ratings_df that combines the
  two DataFrames
9 movie_ids_with_avg_ratings_df =
  ratingsDF.groupBy('movieId').agg(F.count(ratingsDF.rating).alias("count"),
  F.avg(ratingsDF.rating).alias("average"))
10 print 'movie_ids_with_avg_ratings_df:'
11 movie_ids_with_avg_ratings_df.show(3, truncate=False)
12
13 # Note: movie_names_df is a temporary variable, used only to separate the
  steps necessary
14 # to create the movie_names_with_avg_ratings_df DataFrame.
15 #movie_names_df = movie_ids_with_avg_ratings_df.<FILL_IN> # use
  df.join(df2, df.name == df2.name)
16 movie_names_df = movie_ids_with_avg_ratings_df.join(movies_DF,
  movie_ids_with_avg_ratings_df.movieId == movies_DF.ID)
17
18 #movie_names_with_avg_ratings_df = movie_names_df.<FILL_IN>
19 movie_names_with_avg_ratings_df = movie_names_df.select("average", "title",
  "count", "movieId")
20
21 print 'movie_names_with_avg_ratings_df:'
22 movie_names_with_avg_ratings_df.show(3, truncate=False)

```

### ► (3) Spark Jobs

movie\_ids\_with\_avg\_ratings\_df:

```

+-----+-----+-----+
|movieId|count|average|
+-----+-----+-----+
|3997   |2047 |2.0703468490473864|
|1580   |35580|3.55831928049466 |
|3918   |1246 |2.918940609951846 |
+-----+-----+-----+

```

only showing top 3 rows

movie\_names\_with\_avg\_ratings\_df:

```

+-----+-----+-----+-----+
|average|title|count|movieId|
+-----+-----+-----+-----+
|2.0703468490473864|Dungeons & Dragons (2000)|2047|3997|

```

```
|3.55831928049466 |Men in Black (a.k.a. MIB) (1997)|35580|1580 |
|2.918940609951846 |Hellbound: Hellraiser II (1988) |1246 |3918 |
+-----+-----+-----+-----+
only showing top 3 rows
```

Command took 26.81 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:30:32 AM on

movies\_project  
Cmd 22

```
1 #We want to filter our movies with high ratings but greater than or equal
  to 500 reviews.
2 #movies_with_500_ratings_or_more = movie_names_with_avg_ratings_df.
  <FILL_IN>
3 movies_with_500_ratings_or_more = (movie_names_with_avg_ratings_df
4                                   .where("count >= 500"))
5
6                                   .orderBy(movie_names_with_avg_ratings_df.average.desc()))
7 print 'Movies with highest ratings:'
8 movies_with_500_ratings_or_more.show(20, truncate=False)
```

## ► (2) Spark Jobs

Movies with highest ratings:

```
+-----+-----+-----+-----+
-----+-----+-----+
|average          |title
                  |count|movieId|
+-----+-----+-----+-----+
-----+-----+-----+
|4.446990499637029 |Shawshank Redemption, The (1994)
                  |63366|318    |
|4.364732196832306 |Godfather, The (1972)
                  |41355|858    |
|4.334372207803259 |Usual Suspects, The (1995)
                  |47006|50     |
|4.310175010988133 |Schindler's List (1993)
                  |50054|527    |
|4.275640557704942 |Godfather: Part II, The (1974)
                  |27398|1221   |
|4.2741796572216   |Seven Samurai (Shichinin no samurai) (1954)
                  |11611|2019   |
|4.271333600779414 |Rear Window (1954)
                  |17449|904    |
|4.263182346109176 |Band of Brothers (2001)
                  |4305 |7502   |
|4.258326830670664 |Casablanca (1942)
                  |24349|912    |
```

```

|4.256934865900383 |Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)
|6525 |922 |
|4.24807897901911 |One Flew Over the Cuckoo's Nest (1975)
|29932|1193 |
|4.247286821705426 |Dr. Strangelove or: How I Learned to Stop Worrying and Love
the Bomb (1964)|23220|750 |
|4.246001523229246 |Third Man, The (1949)
|6565 |1212 |
|4.235410064157069 |City of God (Cidade de Deus) (2002)
|12937|6016 |
|4.2347902097902095|Lives of Others, The (Das Leben der Anderen) (2006)
|5720 |44555 |
|4.233538107122288 |North by Northwest (1959)
|15627|908 |
|4.2326233183856505|Paths of Glory (1957)
|3568 |1178 |
|4.227123123722136 |Fight Club (1999)
|40106|2959 |
|4.224281931146873 |Double Indemnity (1944)
|4909 |3435 |
|4.224137931034483 |12 Angry Men (1957)
|12934|1203 |

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

only showing top 20 rows

Command took 13.62 seconds -- by meghana.rwgsq1@gmail.com at 8/31/2017, 10:33:40 AM on

movies\_project  
cmd 23

```

1 | #Find all the movies that have average ratings greater than 4.0.
2 | #ratingsDF.groupBy("movieId").avg("rating").filter("avg(rating) >
   | 4.0").show()

```

## ► (2) Spark Jobs

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|movieId|      avg(rating)|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 89056|          4.125|
| 78064|4.111111111111111|
| 81501|          4.5|
| 102119|          4.5|
| 97092|          4.5|
| 858|4.364732196832306|
| 48780|4.042195403318839|
| 118258|          4.5|
| 3226|4.666666666666667|

```

104317	5.0
97872	4.5
3000	4.096298619824341
1303	4.040199611865816
80906	4.005541346973572
116183	4.5
3089	4.143596986817326
104829	4.083333333333333
56490	4.1
125599	5.0
1223	4.066765197275415

only showing top 20 rows

Command took 10.31 seconds -- by meghana.rwgsq@gmail.com at 7/4/2017, 5:44:27 PM on

Movies\_Project  
Cmd 24

```
1 #Find all the movies that have average ratings greater than 4.0 and have
  recieved more than 400 reviews
2
3 ratingsDF.groupBy("movieId")\
4 .agg({"rating":"avg", "*":"count"})\
5 .filter("avg(rating) > 4.0 AND count(1) > 400")\
6 .show()
```

### ► (3) Spark Jobs

movieId	avg(rating)	count(1)
858	4.364732196832306	41355
48780	4.042195403318839	11269
3000	4.096298619824341	9564
1303	4.040199611865816	3607
80906	4.005541346973572	1173
3089	4.143596986817326	3186
1223	4.066765197275415	7781
898	4.171426401336777	6583
5995	4.053922967189729	10515
296	4.174231169217055	67310
68954	4.038266407599309	9264
86377	4.08361391694725	891
58559	4.220129171151776	20438
593	4.17705650958151	63299
1199	4.029590886293616	13957
1212	4.246001523229246	6565
7132	4.047279792746114	1544

950	4.184187016081	3358
1198	4.219009123455364	43295
7587	4.065116279069767	645

only showing top 20 rows

Command took 14.55 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:34:20 AM on

movies\_project  
Cmd 25

```
1 #Find all the movies that have less than 10 reviews
2
3 ratingsDF.groupBy("movieId")\
4 .agg({"*":"count"})\
5 .filter("count(1) <10")\
6 .show()
```

#### ► (1) Spark Jobs

movieId	count(1)
104064	5
96469	7
104508	2
46952	9
102798	8
80451	9
7754	4
119432	4
89056	4
89844	6
92182	2
78064	9
80033	3
83250	9
81349	2
103747	3
71995	5
86927	3
53963	5
69042	4

only showing top 20 rows

Command took 9.93 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 9:32:12 PM on

movies\_project  
Cmd 26

```

1 # Collaborative filtering
2 # Splitting Data into Train and test
3 (training, test) = ratingsDF.randomSplit([0.8, 0.2])

```

Command took 0.04 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:39:04 AM on

movies\_project  
Cmd 27

```
1 training.show()
```

#### ► (1) Spark Jobs

userId	movieId	rating	timestamp
1	2	3.5	2005-04-02 23:53:...
1	29	3.5	2005-04-02 23:31:...
1	47	3.5	2005-04-02 23:32:...
1	50	3.5	2005-04-02 23:29:...
1	112	3.5	2004-09-10 03:09:...
1	151	4.0	2004-09-10 03:08:...
1	253	4.0	2005-04-02 23:35:...
1	260	4.0	2005-04-02 23:33:...
1	293	4.0	2005-04-02 23:31:...
1	296	4.0	2005-04-02 23:32:...
1	318	4.0	2005-04-02 23:33:...
1	367	3.5	2005-04-02 23:53:...
1	541	4.0	2005-04-02 23:30:...
1	593	3.5	2005-04-02 23:31:...
1	653	3.0	2004-09-10 03:08:...
1	919	3.5	2004-09-10 03:07:...
1	924	3.5	2004-09-10 03:06:...
1	1009	3.5	2005-04-02 23:53:...
1	1036	4.0	2005-04-02 23:44:...
1	1079	4.0	2004-09-10 03:07:...

only showing top 20 rows

Command took 15.41 seconds -- by meghana.rwgsq@gmail.com at 7/4/2017, 5:27:31 PM on

Movies\_Project  
Cmd 28

```

1 #comparing Training and Test Dataset
2 tempDF = training.join(test, training.movieId == test.movieId)
3 display(tempDF)

```

#### ► (1) Spark Jobs



userId	movieId	rating	timestamp
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000
603	148	2	1996-07-25T15:28:36.000+0000

Showing the first 1000 rows.



Command took 50.24 seconds -- by meghana.rwgsq@gmail.com at 6/29/2017, 7:36:38 PM on

Movies\_project  
Cmd 29

```

1 #Dataset in Training and not in Test
2 filtered_test = training.subtract(test)
3
4 display (filtered_test)
5 filtered_test.count()

```

## ► (2) Spark Jobs

userId	movieId	rating
13	509	4
14	31658	4
21	953	3
21	1282	3
21	4034	4
24	2539	2
24	4226	4
26	587	4
26	300	2

Showing the first 1000 rows.



Command took 1.90 minutes -- by meghana.rwgsq@gmail.com at 6/29/2017, 8:08:32 PM on

Movies\_project  
Cmd 30

```

1 print "Training set size: ", training.count()
2 print "Test set size: ", test.count()
3 #print "Validation set size: ", test.count()

```

#### ► (2) Spark Jobs

Training set size: 15998929

Test set size: 4001334

Command took 1.13 minutes -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:39:13 AM on

movies\_project  
Cmd 31

```

1 # Build the recommendation model using ALS on the training data
2 # Note we set cold start strategy to 'drop' to ensure we don't get NaN
  evaluation metrics
3 from pyspark.sql import SparkSession
4 from pyspark.ml.evaluation import RegressionEvaluator
5 from pyspark.ml.recommendation import ALS
6 from pyspark.sql import Row
7
8 als = ALS(rank=10, maxIter=10, regParam=0.01, userCol="userId",
  itemCol="movieId", ratingCol="rating")
9 #als = ALS(maxIter=5, regParam=0.01, userCol="userId", itemCol="movieId",
  ratingCol="rating", coldStartStrategy="drop")
10 model = als.fit(training)

```

#### ► (5) Spark Jobs

Command took 2.11 minutes -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:40:44 AM on

movies\_project  
Cmd 32

```

1 #predictions = model.transform(test).dropna()
2 predictions = model.transform(test)
3 predictions = predictions.dropna()
4 predictions.registerTempTable("predictions_table")

```

Command took 0.12 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:42:57 AM on

movies\_project  
Cmd 33

```

1 %sql select user_id, movie_id, rating, prediction from predictions
2 spark.sql("select userId, movieId, rating, prediction from
  predictions_table").show()
3

```

#### ► (1) Spark Jobs

```

+-----+-----+-----+-----+
|userId|movieId|rating|prediction|

```

```

+-----+-----+-----+-----+
| 92852|    148|    3.0| 2.5335486|
| 81218|    148|    1.0| 2.7732363|
| 91782|    148|    3.0| 2.9888206|
| 13170|    148|    3.0| 0.0466154|
|  1259|    148|    5.0| 3.2165732|
| 44882|    148|    4.0| 2.3504906|
| 94994|    148|    4.0|  3.119975|
| 90757|    148|    3.0| 2.9625764|
|  3673|    148|    2.0| 2.6216304|
| 64843|    148|    3.5|  2.630504|
| 81300|    148|    1.0| 2.3215954|
|118205|    148|    3.5| 2.8588994|
|109121|    148|    4.0|  3.454861|
| 68360|    148|    2.0|  2.942938|
|132268|    148|    2.0| 1.5295749|
| 30699|    148|    3.0| 2.1022847|
| 28361|    148|    4.0|  4.177149|
| 36723|    148|    1.0| 1.9845809|
|  9084|    148|    2.0| 3.0726907|
| 66709|    148|    5.0|  3.728733|
+-----+-----+-----+-----+

```

only showing top 20 rows

Command took 35.43 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 10:43:05 AM on

movi3s\_project  
Cmd 34

```

1 #Generate top 10 movie recommendations for each user
2 userRecs = model.recommendForAllUsers(10)
3 userRecs.show(10, truncate = False)

```

#### ► (1) Spark Jobs

```

+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|userId|recommendations
|
+-----+-----+-----+-----+
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|148   |[[128366,10.951867], [1830,10.697975], [40969,9.900242], [27108,9.73136
5], [76022,9.663902], [69310,9.3615885], [94074,9.253457], [80189,8.878327], [73
170,8.857643], [87719,8.633941]]
|463   |[[116155,7.8540316], [87884,7.6012306], [76022,7.536461], [69931,7.45520
8], [107780,7.1698666], [69666,7.0395384], [72045,6.9239545], [87040,6.9138904],

```

```
[117308,6.8977456], [109633,6.892449]]|
|471  |[[74159,9.095757], [116951,8.812828], [76022,8.2325325], [27108,7.477155
7], [69310,7.318554], [88092,7.197364], [73501,6.787854], [61913,6.649144], [921
69,6.524055], [39244,6.459038]]|
|496  |[[73529,7.6900644], [87884,7.430202], [69464,7.1698527], [88092,7.154989
2], [72388,7.055085], [116951,6.789952], [86572,6.6127315], [8411,6.582972], [84
796,6.551797], [73139,6.5232415]]|
|833  |[[116951,13.432829], [61913,11.178102], [98126,10.112643], [74159,9.9142
13], [27108,9.62832], [104861,9.517236], [81075,9.331002], [26027,9.306303], [33
132,9.176152], [81059,9.121372]]|
|1088 |[[86947,10.466293], [72388,10.142122], [92696,8.911163], [110380,8.63503
4], [104390,8.579515], [97715,8.305207], [44861,8.267274], [629,8.1867485], [777
19,8.138725], [55856,8.128115]]|
|1238 |[[34729,7.121804], [27108,6.643126], [69931,6.5108547], [88092,6.18079
3], [73879,6.144084], [73529,6.084306], [67894,6.067928], [86947,6.0549393], [11
0380,6.03752], [82304,6.0205054]]|
|1342 |[[6674,8.4584055], [27679,8.04373], [68976,7.850332], [72360,7.8449745],
[98595,7.82905], [72735,7.717164], [92169,7.561826], [8954,7.550848], [99387,7.5
01115], [57478,7.414889]]|
|1580 |[[1830,16.55101], [128366,13.99382], [73170,13.127473], [40969,11.811542
5], [53548,11.645727], [116155,11.541295], [80189,11.204805], [92314,10.934051
5], [3209,10.876148], [82061,10.749665]]|
|1591 |[[53476,9.660441], [39244,9.43916], [87040,9.315279], [74159,9.262892],
[66579,9.137938], [59302,8.514868], [76022,8.449947], [82261,8.342544], [82842,
8.307518], [77433,8.128576]]|
```

```
+-----+-----+
-----+
-----+
only showing top 10 rows
```

Command took 56.02 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:03:41 AM on

movies\_project  
cmd 35

```
1 # Generate top 10 user recommendations for each movie
2 movieRecs = model.recommendForAllItems(10)
3 movieRecs.show(10,truncate = False)
```

#### ► (1) Spark Jobs

```
+-----+-----+
-----+
-----+
|movieId|recommendations|
|
+-----+-----+
-----+
```

```

-----+
|148      |[[33314,8.701766], [1425,8.180933], [81312,7.373311], [68156,7.208327
3], [106990,6.8708425], [86490,6.8576126], [131667,6.6769357], [113811,6.57130
2], [42589,6.48554], [36374,6.4662094]] |
|463      |[[120759,7.827287], [103795,7.3661604], [1425,7.3247886], [7637,7.02861
55], [41182,6.988009], [65884,6.946594], [46236,6.7009373], [101608,6.600505],
[117304,6.5729437], [130899,6.54625]] |
|471      |[[15629,7.3206534], [39244,7.006679], [68156,6.9869514], [55809,6.80674
6], [35429,6.7478523], [17999,6.739544], [107667,6.633784], [132964,6.6239376],
[112794,6.5889277], [54192,6.460428]] |
|496      |[[1425,11.919976], [33314,8.572311], [42589,8.134817], [120759,7.643493
7], [101255,7.342348], [26714,7.2738714], [86490,7.260195], [59542,7.174942], [7
2203,7.1624765], [36374,7.149855]] |
|833      |[[7245,6.9702134], [138324,6.848807], [86490,6.721907], [121230,6.65824
5], [126728,6.6005583], [29030,6.476948], [29388,6.4507437], [114637,6.3349886],
[51105,6.30784], [67679,6.2839894]] |
|1088     |[[61007,9.486544], [92390,8.747037], [7332,7.796366], [101608,7.43559
2], [95345,7.182036], [99906,7.117282], [8892,6.961412], [65884,6.9582014], [132
685,6.84749], [116845,6.832096]] |
|1238     |[[61315,7.817972], [81312,6.8869853], [101255,6.847694], [110290,6.7657
6], [14403,6.6632776], [105116,6.627554], [27841,6.5338273], [7326,6.5334315],
[92724,6.456632], [73969,6.453871]] |
|1342     |[[1425,7.136017], [125248,6.5628214], [97609,6.553137], [122900,6.55167
8], [49826,6.5504475], [11941,6.4981246], [22946,6.486262], [124859,6.455388],
[9635,6.372327], [9340,6.367054]] |
|1580     |[[36089,6.003181], [861,5.9432993], [112095,5.7447734], [46278,5.713271
6], [127916,5.7111936], [71611,5.683817], [117250,5.68226], [115837,5.658524],
[52622,5.6481333], [22718,5.6437683]] |
|1591     |[[116848,6.179776], [1425,6.070424], [35823,6.035642], [69716,5.94612
4], [96611,5.924816], [95981,5.9094105], [86490,5.883312], [112156,5.8413234],
[28134,5.7672577], [27735,5.762802]] |

```

```

+-----+-----+
-----+
-----+

```

only showing top 10 rows

Command took 51.01 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:02:35 AM on

movies\_project  
Cmd 36

```

1 | spark.sql("select count(userId) as Users, movieId from predictions_table
   | group by movieId LIMIT 20").show()

```

#### ► (1) Spark Jobs

```

+-----+-----+
|Users|movieId|
+-----+-----+

```

	67	148
	76	463
	2253	471
	80	496
	283	833
	2161	1088
	598	1238
	653	1342
	7026	1580
	1128	1591
	2308	1645
	52	1829
	1023	1959
	498	2122
	424	2142
	1291	2366
	42	2659
	258	2866
	2697	3175
	22	3749
+-----+-----+		

Command took 38.77 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 9:41:37 PM on  
movies\_project  
Cmd 37

```

1 from pyspark.sql import functions as F
2
3 # Rename column from movieId to ID
4 movies_DF = moviesDF.withColumnRenamed('movieId', 'ID')
5
6 # From predictions, create a movie_ids_with_avg_ratings_df that combines
  the two DataFrames
7 movie_ids_with_avg_predicted_ratings_df =
  predictions.groupBy('movieId').agg(F.count(predictions.prediction).alias("c
    ount"), F.avg(predictions.prediction).alias("average"))
8 #print 'movie_ids_with_avg_predicted_ratings_df:'
9 #movie_ids_with_avg_predicted_ratings_df.show(3, truncate=False)
10
11 # Note: movie_names_df is a temporary variable, used only to separate the
  steps necessary
12 # to create the movie_names_with_avg_ratings_df DataFrame.
13 #movie_names_df = movie_ids_with_avg_ratings_df.<FILL_IN> # use
  df.join(df2, df.name == df2.name)
14 predicted_movie_names_df =
  movie_ids_with_avg_predicted_ratings_df.join(movies_DF,
  movie_ids_with_avg_predicted_ratings_df.movieId == movies_DF.ID)
15
16 #movie_names_with_avg_ratings_df = movie_names_df.<FILL_IN>
17 movie_names_with_avg_pred_ratings_df =
  predicted_movie_names_df.select("average", "title", "count", "movieId")
18
19 #print 'movie_names_with_avg_ratings_df:'
20 #movie_names_with_avg_pred_ratings_df.show(3, truncate=False)
21
22 #We want to filter our movies with high ratings but greater than or equal
  to 500 reviews.
23 #movies_with_500_ratings_or_more = movie_names_with_avg_ratings_df.
  <FILL_IN>
24 predicted_movies_with_500_ratings_or_more =
  (movie_names_with_avg_pred_ratings_df
25     .where("count >= 500")
26     .orderBy(movie_names_with_avg_pred_ratings_df.average.desc()))
27
28 print 'Movies with highest ratings:'
29 predicted_movies_with_500_ratings_or_more.show(20, truncate=False)

```

- ▶ (2) Spark Jobs

```
movie_ids_with_avg_predicted_ratings_df:
```

Movies with highest ratings:

---

```

-----+-----+-----+
|average          |title
                  |count|movieId|
-----+-----+-----+
|4.276359616133422 |Shawshank Redemption, The (1994)
                  |12513|318    |
|4.240856685485456 |Third Man, The (1949)
                  |1317 |1212   |
|4.219244047411683 |Godfather, The (1972)
                  |8345 |858    |
|4.193616895280234 |Usual Suspects, The (1995)
                  |9424 |50     |
|4.185131523330215 |Seven Samurai (Shichinin no samurai) (1954)
                  |2359 |2019   |
|4.183755558500053 |Rear Window (1954)
                  |3538 |904    |
|4.181936215899116 |Casablanca (1942)
                  |4949 |912    |
|4.180760424454526 |Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)
                  |1322 |922    |
|4.164261522881997 |Dr. Strangelove or: How I Learned to Stop Worrying and Love
the Bomb (1964)|4591 |750    |
|4.159615292880716 |Schindler's List (1993)
                  |9804 |527    |
|4.157970987149139 |Paths of Glory (1957)
                  |737  |1178   |
|4.151550216978014 |North by Northwest (1959)
                  |3161 |908    |
|4.147988171927924 |Notorious (1946)
                  |966  |930    |
|4.147914183192143 |Godfather: Part II, The (1974)
                  |5466 |1221   |
|4.147455377140265 |Big Sleep, The (1946)
                  |1131 |1284   |
|4.145923303171358 |Double Indemnity (1944)
                  |963  |3435   |
|4.145621530264448 |12 Angry Men (1957)
                  |2649 |1203   |
|4.141916454161409 |Rashomon (Rashômon) (1950)
                  |738  |5291   |
|4.137867432680226 |Band of Brothers (2001)
                  |899  |7502   |
|4.1342800594020535|Touch of Evil (1958)
                  |925  |1248   |
-----+-----+-----+
-----+-----+-----+

```



only showing top 20 rows

Command took 39.79 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 11:13:38 AM on

movies\_project  
Cmd 38

```
1 spark.sql("select count(*) from predictions_table").show()
```

► (1) Spark Jobs


```
+-----+
|count(1)|
+-----+
| 4001283|
+-----+
```

Command took 39.19 seconds -- by meghana.rwgsql@gmail.com at 8/30/2017, 4:23:28 PM on

Movigs\_Recommendation  
Cmd 39

```
1 display (spark.sql("select rating, prediction from predictions_table"))
```

► (1) Spark Jobs



Showing sample based on the first 1000 rows.



Command took 35.01 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 11:16:26 AM on


movies\_project  
Cmd 40

```

1
2 display(predictions.sample(False, 1000.0/ratingsDF.count()))

```

► (4) Spark Jobs



Command took 46.74 seconds -- by meghana.rwgsq@gmail.com at 8/30/2017, 9:59:47 PM on

movies\_project  
Cmd 41

```

1 # Evaluate the model by computing the RMSE on the test data
2 evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
3                               predictionCol="prediction")
4 rmse = evaluator.evaluate(predictions)
5 print("Root-mean-square error = " + str(rmse))

```

► (1) Spark Jobs

Root-mean-square error = 0.804667235236

Command took 39.97 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:17:09 AM on

movies\_project  
Cmd 42

```

1 displayHTML("<h4>The Root-mean-square error is %s</h4>" % str(rmse))

```

## The Root-mean-square error is 0.804667235236

Command took 0.04 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:17:59 AM on

movies\_project  
Cmd 43

```

1 Movies_recommended = spark.sql("SELECT predictions_table.movieId,
  movies.title, rating, prediction \
2                               FROM predictions_table JOIN movies on
  movies.movieId = predictions_table.movieId \
3                               Group by predictions_table.movieId,
  movies.title, rating, prediction \
4                               Order by prediction DESC, rating DESC LIMIT
  20")
5 Movies_recommended.show(truncate = False)

```

► (2) Spark Jobs

movieId	title	rating	prediction
3455	Buddy Boy (1999)	1.0	9.837434
97059	Katy Perry: Part of Me (2012)	0.5	9.402147
99861	Jesse Stone: Sea Change (2007)	5.0	9.29036
26136	Hallelujah Trail, The (1965)	2.5	9.207833
119145	Kingsman: The Secret Service (2015)	1.0	8.797971
1165	Bloody Child, The (1996)	5.0	8.539613
46848	Gumball Rally, The (1976)	4.5	8.5312395
39416	Kids in America (2005)	5.0	8.268056
2710	Blair Witch Project, The (1999)	3.0	8.222418
63992	Twilight (2008)	5.0	8.204944
78772	Twilight Saga: Eclipse, The (2010)	5.0	8.2048855
71141	Airbag (1997)	0.5	8.130741
501	Naked (1993)	5.0	8.091531
1354	Breaking the Waves (1996)	5.0	8.080061
4491	Criminal Law (1988)	0.5	8.074579
80505	2012: Supernova (2009)	4.0	8.061749
7318	Passion of the Christ, The (2004)	5.0	8.057657
5922	Attila (Attila Flagello di Dio) (1982)	4.5	8.050166
67501	Kogel mogel (1988)	3.0	8.045789
26339	Dolemite (1975)	5.0	8.014356

Command took 44.85 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:18:10 AM on

movies\_project  
Cmd 44

```

1 #Recommended movies
2 Movies_Recommended_from_model = spark.sql("SELECT
   predictions_table.movieId, movies.title, avg(prediction) as
   Predicted_ratings \
3                                     FROM predictions_table \
4                                     JOIN movies ON
   movies.movieId=predictions_table.movieId \
5                                     GROUP BY
   predictions_table.movieId, movies.title \
6                                     ORDER BY AVG(prediction) DESC\
7                                     LIMIT 20")
8 Movies_Recommended_from_model.show(truncate = False)

```

## ► (2) Spark Jobs

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|movieId|title                                     |Predicted_ratings |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|107357 |Call Me Crazy: A Five Film (2013)         |7.227933883666992 |
|56022  |Harrison Bergeron (1995)                   |6.917505264282227 |
|55659  |Return to the 36th Chamber (Shao Lin da peng da shi) (1980)|6.843296051025391 |
|97196  |Shottas (2002)                            |6.524924039840698 |
|80324  |Abandoned (2010)                          |6.486706256866455 |
|89961  |Play (2011)                               |6.443424701690674 |
|95021  |Outer Space (2000)                        |6.418587684631348 |
|98834  |Fitzgerald Family Christmas, The (2012)   |6.3923845291137695|
|98755  |1911 (Xinhai geming) (2011)               |6.2648186683654785|
|78332  |Beautiful Person, The (La belle personne) (2008)|6.190258026123047 |
|109416 |Bring It On: Fight to the Finish (2009)   |6.165205478668213 |
|96473  |Prime Suspect: Inner Circles (1995)       |6.139786720275879 |
|41650  |Mother India (1957)                       |6.084947109222412 |

```

```

|49872  |Loose Change: Second Edition (2006)
                                     |5.997299671173096 |
|74937  |Two-Minute Warning (1976)
                                     |5.947014331817627 |
|84506  |Silent Souls (Ovsyanki) (2010)
                                     |5.800346374511719 |
|99861  |Jesse Stone: Sea Change (2007)
                                     |5.757004082202911 |
|27372  |Uprising (2001)
                                     |5.701282978057861 |
|7441   |Thousand Clouds of Peace, A (Mil nubes de paz cercan el cielo, amor, ja
más acabarás de ser amor) (2003)|5.665889501571655 |
|81906  |Snow and Ashes (2010)
                                     |5.659394264221191 |
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+

```

Command took 41.43 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:19:31 AM on

movies\_project  
Cmd 45

```
1 predictions.describe().show()
```

#### ► (1) Spark Jobs

```

+-----+-----+-----+-----+-----+
---+
|summary|      userId|      movieId|      rating|      predict
ion|
+-----+-----+-----+-----+-----+
---+
|  count|      4000316|      4000316|      4000316|      4000
316|
|   mean|69017.89363490284|  9013.73503618214|  3.525570604922211|3.4412456563136
664|
| stddev|40043.73561717234|19720.687189242326|1.0519828325185538|0.7479698634538
245|
|   min|           1|           1|           0.5|      -6.137
381|
|   max|      138493|      131013|           5.0|           9.837
434|
+-----+-----+-----+-----+-----+
---+

```

Command took 38.73 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:20:32 AM on

movies\_project  
Cmd 46

```
1 moviesRDD=moviesDF.rdd
2 ratingsRDD=ratingsDF.rdd
```

Command took 0.12 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:47:08 PM on  
movies\_project  
Cmd 47

```
1 type(moviesRDD)
2 type(ratingsRDD)
```

Out[50]: pyspark.rdd.RDD

Command took 0.07 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:47:11 PM on  
movies\_project  
Cmd 48

```
1 from pyspark import SparkConf, SparkContext
2
3 numRatings = ratingsRDD.count()
4 numUsers   = ratingsRDD.map(lambda r:r[0]).distinct().count()
5 numMovies  = ratingsRDD.map(lambda r:r[1]).distinct().count()
6 print "--- %d ratings from %d users for %d movies\n" % (numRatings,
  numUsers, numMovies)
```

### ► (3) Spark Jobs

--- 20000263 ratings from 138493 users for 26744 movies

Command took 9.94 minutes -- by meghana.rwgsq@gmail.com at 8/30/2017, 5:34:53 PM on  
Movies\_Recommendation  
Cmd 49

```

1 #Adding new user ratings.
2 #Now we need to rate some movies for the new user. We will put them in a
  new RDD and we will use the user ID 0, that is not assigned in the
  MovieLens dataset
3
4 new_user_ID = 0
5 # The format of each line is (userID, movieID, rating)
6 new_user_ratings = [
7 (0,318,4.5,1112484580), # Shawshank Redemption (1994)
8 (0,858,4.7,1112484940), # Godfather (1972)
9 (0,50,5.0,1094785709), # Usual suspects (1995)
10 (0,527,5.0,1094785691), # Schindlers List (1993)
11 (0,1221,5.0,1094785759), # Godfather: Part II (1974)
12 (0,2019,4.8,1112484735), # Seven Samurai (1954)
13 (0,904,4.5,1094786062), # Rear window (1954)
14 (0,7502,4.7,1094785764), # Band of Brothers (2001)
15 (0,912,4.2,1112486150) , # Casablanca (1942)
16 (0,922,4.0,1112486098) # Sunset blvd (1950)
17 ]
18 new_user_ratings_RDD = sc.parallelize(new_user_ratings)
19 print 'New user ratings: %s' % new_user_ratings_RDD.take(10)
20

```

### ► (3) Spark Jobs

New user ratings: [(0, 318, 4.5, 1112484580), (0, 858, 4.7, 1112484940), (0, 50, 5.0, 1094785709), (0, 527, 5.0, 1094785691), (0, 1221, 5.0, 1094785759), (0, 2019, 4.8, 1112484735), (0, 904, 4.5, 1094786062), (0, 7502, 4.7, 1094785764), (0, 912, 4.2, 1112486150), (0, 922, 4.0, 1112486098)]

Command took 0.47 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:47:23 PM on  
 movies\_project  
 Cmd 50

```

1 new_user_ratings =
  new_user_ratings_RDD.toDF(['userId','movieId','rating','timestamp'])

```

### ► (1) Spark Jobs

Command took 0.12 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:49:31 PM on  
 movies\_project  
 Cmd 51

```

1 new_user_ratings = new_user_ratings.withColumn('rating',
  new_user_ratings["rating"].cast("float"))
2 new_user_ratings = new_user_ratings.withColumn('userId',
  new_user_ratings["userId"].cast("int"))
3 new_user_ratings = new_user_ratings.withColumn('movieId',
  new_user_ratings["movieId"].cast("int"))
4 new_user_ratings.show()

```

### ► (3) Spark Jobs

```

+-----+-----+-----+-----+
|userId|movieId|rating| timestamp|
+-----+-----+-----+-----+
|      0|      318|    4.5|1112484580|
|      0|      858|    4.7|1112484940|
|      0|       50|    5.0|1094785709|
|      0|      527|    5.0|1094785691|
|      0|     1221|    5.0|1094785759|
|      0|     2019|    4.8|1112484735|
|      0|      904|    4.5|1094786062|
|      0|     7502|    4.7|1094785764|
|      0|      912|    4.2|1112486150|
|      0|      922|    4.0|1112486098|
+-----+-----+-----+-----+

```

Command took 0.22 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:49:35 PM on  
 movies\_project  
 Cmd 52

```

1 from pyspark.sql.functions import from_unixtime
2 new_user_ratings = new_user_ratings.withColumn('timestamp',
  from_unixtime(new_user_ratings.timestamp))

```

Command took 0.04 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:49:44 PM on  
 movies\_project  
 Cmd 53

```

1 from pyspark.sql.types import TimestampType
2 new_user_ratings = new_user_ratings.withColumn('timestamp',
  new_user_ratings.timestamp.cast(TimestampType()))

```

Command took 0.07 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:49:47 PM on  
 movies\_project  
 Cmd 54

```

1 new_user_ratings.printSchema()

```

root

```
|-- userId: integer (nullable = true)
```



```
|-- movieId: integer (nullable = true)
|-- rating: float (nullable = true)
|-- timestamp: timestamp (nullable = true)
```

Command took 0.04 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:49:51 PM on  
movies\_project  
Cmd 55

```
1 training_with_my_ratings = training.unionAll(new_user_ratings)
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:50:58 PM on  
movies\_project  
Cmd 56

```
1 display(new_user_ratings.limit(10))
```

### ► (3) Spark Jobs

userId	movieId	rating
0	353	4.7
0	50	5
0	527	5
0	1221	5
0	2019	4.8
0	904	4.5
0	7502	4.7
0	912	4.2
0	922	4



Command took 0.22 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:24:53 PM on  
movies\_project  
Cmd 57

```
1 # TODO: Replace <FILL IN> with appropriate code
2 training_with_my_ratings = training.unionAll(new_user_ratings)
3
4 print ('The training dataset now has %s more entries than the original
   training dataset' %
5       (training_with_my_ratings.count() - training.count()))
6 assert (training_with_my_ratings.count() - training.count()) ==
   new_user_ratings.count()
7 #The training dataset now has 11 more entries than the original training
   dataset
```

### ► (5) Spark Jobs

The training dataset now has 10 more entries than the original training dataset

Command took 1.96 minutes -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:21:36 PM on

```
movies_project
Cmd 58
1 #Training new dataset with ALS model
2
3 # Build the new recommendation model using ALS on the training data
4 # Note we set cold start strategy to 'drop' to ensure we don't get NaN
  evaluation metrics
5 from pyspark.sql import SparkSession
6 from pyspark.ml.evaluation import RegressionEvaluator
7 from pyspark.ml.recommendation import ALS
8 #from pyspark.mllib.recommendation import ALS, Rating
9 from pyspark.sql import Row
10
11 als = ALS(rank=10, maxIter=10, regParam=0.01, userCol="userId",
  itemCol="movieId", ratingCol="rating")
12 #als = ALS(maxIter=5, regParam=0.01, userCol="userId", itemCol="movieId",
  ratingCol="rating", coldStartStrategy="drop")
13 new_rating_model = als.fit(training_with_my_ratings)
14
15 #print "New model trained in %s seconds" % round(tt,3)
16
```

### ► (5) Spark Jobs

Command took 2.10 minutes -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:51:16 PM on

```
movies_project
Cmd 59
1 #predictions = model.transform(test).dropna()
2 predictions_with_my_ratings = new_rating_model.transform(test)
3 predictions_with_my_ratings = predictions_with_my_ratings.dropna()
4 predictions_with_my_ratings.registerTempTable("predictions_with_my_ratings_
  table")
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:28:27 PM on

```
movies_project
Cmd 60
1 spark.sql("select * from predictions_with_my_ratings_table order by
  prediction desc LIMIT 10").show()
2
```

### ► (1) Spark Jobs

```

+-----+-----+-----+-----+-----+
|userId|movieId|rating|          timestamp|prediction|
+-----+-----+-----+-----+-----+
| 24994| 100010|  4.0|2015-02-20 22:19:47| 10.508722|
|138215|   876|  5.0|2002-11-10 06:20:00| 10.221236|
| 51007|  3228|  2.0|2000-08-28 20:06:50|  9.273807|
| 47938|  67501|  3.0|2010-02-06 19:16:39|  9.200949|
|119833| 111320|  5.0|2014-07-31 03:24:01|  8.994706|
|117621|  4244|  5.0|2007-03-28 09:35:20|  8.776396|
| 78022|  7441|  3.5|2004-08-10 03:14:23|  8.760872|
| 24480|  5922|  4.5|2006-07-12 03:20:59|  8.693072|
|104117|  56093|  5.0|2010-02-07 09:21:14|  8.650107|
|113325|  44595|  0.5|2010-12-17 07:36:11|  8.552686|
+-----+-----+-----+-----+-----+

```

Command took 37.57 seconds -- by meghana.rwgsq1@gmail.com at 8/31/2017, 2:54:54 PM on

movies\_project  
Cmd 61

```

1 # Generate top 10 user recommendations for each movie
2 movieRecsnew = new_rating_model.recommendForAllItems(10)
3 movieRecsnew.show(10,truncate = False)

```

#### ► (1) Spark Jobs

```

+-----+-----+-----+-----+-----+
|movieId|recommendations|
+-----+-----+-----+-----+-----+
|148     |[[124002,8.557087], [89066,8.09482], [43815,7.5858655], [33314,7.496185], [128821,7.2536316], [56145,6.8170147], [112318,6.7748556], [5747,6.746767], [61007,6.746116], [87447,6.7400374]]|
|463     |[[101608,7.5712504], [65884,7.5201693], [92390,7.3081822], [118248,7.1546474], [79298,6.9978046], [7637,6.960828], [28022,6.9503565], [54584,6.9177628], [67339,6.9005127], [112156,6.639238]]|
|471     |[[107804,6.9638305], [68156,6.7373343], [75545,6.385509], [72458,6.2377696], [35429,6.223061], [55716,6.198926], [34369,6.1974516], [53118,6.1938286], [107667,6.1792274], [120483,6.1639795]]|
|496     |[[1425,8.402818], [5625,8.163208], [58845,8.043345], [84580,7.989399], [119988,7.8339767], [17816,7.8166494], [130693,7.7461033], [4273,7.668422], [112469,7.6296096], [74030,7.615759]]|
|833     |[[97184,7.3606925], [75306,6.9260845], [138215,6.9125247], [67679,6.848727], [5474,6.839116], [86166,6.765509], [5747,6.7330513], [121534,6.7165036],

```

```

[121230,6.4955673], [50529,6.48988]]      |
|1088   |[[74576,7.644407], [133404,7.607803], [76772,7.459447], [120759,7.44451
9], [24829,7.4256725], [93821,7.417887], [62342,7.396893], [93809,7.390879], [72
45,7.360171], [110831,7.3552527]]          |
|1238   |[[61315,7.6330214], [53192,6.8270493], [19366,6.5271854], [5747,6.52172
1], [84889,6.5086412], [14403,6.4422936], [14571,6.4190235], [51055,6.383786],
[87447,6.3422785], [9241,6.3329163]]      |
|1342   |[[97609,7.4944735], [124859,6.8256826], [133374,6.8231425], [63539,6.64
2704], [63122,6.6319804], [122900,6.5695662], [114432,6.5227604], [1425,6.513621
3], [68007,6.2603736], [115976,6.259104]] |
|1580   |[[115837,6.137461], [42714,5.9726377], [54192,5.9595323], [20349,5.8790
793], [72725,5.8467574], [34369,5.836102], [88922,5.812204], [36089,5.791688],
[44957,5.7126417], [22718,5.70847]]      |
|1591   |[[1425,7.4258857], [138215,6.4615006], [86490,6.347809], [99021,6.15392
3], [112156,6.141174], [97136,6.1349545], [69716,6.087041], [27735,6.070429], [1
07650,5.999553], [54192,5.922649]]      |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
-----+
only showing top 10 rows

```

Command took 58.34 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 2:50:58 PM on

movies\_project  
Cmd 62

```

1 # Evaluate the model by computing the RMSE on the test data
2 evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
3                               predictionCol="prediction")
4 rmse = evaluator.evaluate(predictions_with_my_ratings)
5 print("Root-mean-square error for new ratings = " + str(rmse))*

```

#### ► (1) Spark Jobs

Root-mean-square error for new ratings = 0.805098946653

Command took 38.66 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 1:30:03 PM on

movies\_project  
Cmd 63

```

1 #PREDICT YOUR RATINGS
2 # Create a list of my rated movie IDs
3 myRatedMovieIds = [x[1] for x in new_user_ratings]
4 myRatedMovieIds

```

Out[89]:

```

[Column<userId[1]>,
 Column<movieId[1]>,
 Column<rating[1]>,
 Column<timestamp[1]>]

```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:44:12 PM on

movies\_project  
Cmd 64

```
1 moviesDF = moviesDF.withColumn('movieId', moviesDF["movieId"].cast("int"))
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:52:22 PM on

movies\_project  
Cmd 65

```
1 moviesDF.printSchema()
```

root

```
|-- movieId: integer (nullable = true)
|-- title: string (nullable = true)
|-- genres: string (nullable = true)
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:52:37 PM on

movies\_project  
Cmd 66

```
1 # Filter out the movies I already rated.
2 notRated_df = moviesDF.filter(~
  moviesDF["movieId"].isin(myRated_movie_ids))
```

AnalysisException: u"Can't extract value from userId#5274: need struct type but got int;"

Command took 0.17 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 1:53:43 PM on

movies\_project  
Cmd 67

```
1 # Rename the "ID" column to be "movieId", and add a column with my_user_id
  as "userId".
2 my_unrated_movies_df = notRated_df.selectExpr("ID as
  movieId").withColumn('userId', F.lit(my_user_id))
3
4 # Use my_rating_model to predict ratings for the movies that I did not
  manually rate.
5 raw_predicted_ratings_df = my_ratings_model.transform(my_unrated_movies_df)
6
7 predicted_ratings_df =
  raw_predicted_ratings_df.filter(raw_predicted_ratings_df['prediction'] !=
    float('nan'))
```

Cmd 68

```
1 predicted_with_counts_df =  
  predicted_ratings_df.join(movie_names_with_avg_ratings_df,movie_names_with_  
    avg_ratings_df["movieId"]==predicted_ratings_df["movieId"])  
2 predicted_highest_rated_movies_df =  
  predicted_with_counts_df.filter(predicted_with_counts_df["count"]>75).sort(  
    "prediction",ascending=False)  
3  
4 print ('My 25 highest rated movies as predicted (for movies with more than  
    75 reviews):')  
5 predicted_highest_rated_movies_df.show(25)
```

Cmd 69

```

1 from pyspark.sql import functions as F
2
3 # Rename column from movieId to ID
4 movies_DF = moviesDF.withColumnRenamed('movieId', 'ID')
5
6 # From predictions, create a movie_ids_with_avg_ratings_df that combines
  the two DataFrames
7 new_movie_ids_with_avg_predicted_ratings_df =
  predictions_with_my_ratings.groupBy('movieId').agg(F.count(predictions_with
    _my_ratings.prediction).alias("count"),
    F.avg(predictions_with_my_ratings.prediction).alias("average"))
8 #print 'movie_ids_with_avg_predicted_ratings_df:'
9 #movie_ids_with_avg_predicted_ratings_df.show(3, truncate=False)
10
11 # Note: movie_names_df is a temporary variable, used only to separate the
  steps necessary
12 # to create the movie_names_with_avg_ratings_df DataFrame.
13 #movie_names_df = movie_ids_with_avg_ratings_df.<FILL_IN> # use
  df.join(df2, df.name == df2.name)
14 new_predicted_movie_names_df =
  new_movie_ids_with_avg_predicted_ratings_df.join(movies_DF,
    new_movie_ids_with_avg_predicted_ratings_df.movieId == movies_DF.ID)
15
16 #movie_names_with_avg_ratings_df = movie_names_df.<FILL_IN>
17 new_movie_names_with_avg_pred_ratings_df =
  new_predicted_movie_names_df.select("average", "title", "count", "movieId")
18
19 #print 'movie_names_with_avg_ratings_df:'
20 #movie_names_with_avg_pred_ratings_df.show(3, truncate=False)
21
22 #We want to filter our movies with high ratings but greater than or equal
  to 500 reviews.
23 #movies_with_500_ratings_or_more = movie_names_with_avg_ratings_df.
  <FILL_IN>
24 new_predicted_movies_with_500_ratings_or_more =
  (new_movie_names_with_avg_pred_ratings_df
25     .where("count >= 500")
26
  .orderBy(new_movie_names_with_avg_pred_ratings_df.average.desc()))
27
28 print 'Movies with highest ratings:'
29 new_predicted_movies_with_500_ratings_or_more.show(20, truncate=False)

```

- ▶ (2) Spark Jobs

Movies with highest ratings:

```

-----+-----+-----+
|average          |title
                  |count|movieId|
-----+-----+-----+
|4.254283139357548|Shawshank Redemption, The (1994)
                  |12513|318    |
|4.22503672482484 |Third Man, The (1949)
                  |1317 |1212   |
|4.195024231735141|Godfather, The (1972)
                  |8345 |858    |
|4.175523698708929|Seven Samurai (Shichinin no samurai) (1954)
                  |2359 |2019   |
|4.17406071068002 |Rear Window (1954)
                  |3538 |904    |
|4.173247438199274|Usual Suspects, The (1995)
                  |9424 |50     |
|4.172001171513332|Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)
                  |1322 |922    |
|4.170191933998413|Casablanca (1942)
                  |4949 |912    |
|4.152172995341341|Dr. Strangelove or: How I Learned to Stop Worrying and Love t
he Bomb (1964)|4591 |750    |
|4.140182386042633|Schindler's List (1993)
                  |9804 |527    |
|4.138639154275057|North by Northwest (1959)
                  |3161 |908    |
|4.138031953557074|Rashomon (Rashômon) (1950)
                  |738  |5291   |
|4.135877628541436|Big Sleep, The (1946)
                  |1131 |1284   |
|4.134845475108205|Double Indemnity (1944)
                  |963  |3435   |
|4.131054048945816|Paths of Glory (1957)
                  |737  |1178   |
|4.130510358326207|12 Angry Men (1957)
                  |2649 |1203   |
|4.127649324721303|Thin Man, The (1934)
                  |638  |950    |
|4.126916712084419|Notorious (1946)
                  |966  |930    |
|4.125684980212188|All About Eve (1950)
                  |943  |926    |
|4.125092605320183|Touch of Evil (1958)
                  |925  |1248   |
-----+-----+-----+
-----+-----+-----+

```



only showing top 20 rows

Command took 46.45 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 2:58:27 PM on

movies\_project  
Cmd 70

```
1 display(new_predicted_movies_with_500_ratings_or_more.limit(10))
```

► (2) Spark Jobs

average	title
4.254283139357548	Shawshank Redemption, The (1994)
4.22503672482484	Third Man, The (1949)
4.195024231735141	Godfather, The (1972)
4.175523698708929	Seven Samurai (Shichinin no samurai) (1954)
4.17406071068002	Rear Window (1954)
4.173247438199274	Usual Suspects, The (1995)
4.172001171513332	Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)
4.170191933998413	Casablanca (1942)
4.152172005211241	Dr. Strangelove or: How I Learned to Stop Worrying



Command took 48.04 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 3:58:20 PM on

movies\_project  
Cmd 71

```
1 predictions_with_my_ratings.registerTempTable("predictions_with_my_ratings_
table")
```

Cmd 72

```
1 new_predicted_movies_with_500_ratings_or_more.registerTempTable("Recommende
d_movies")
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 3:03:27 PM on

movies\_project  
Cmd 73

```
1 spark.sql("select * from Recommended_movies limit 10").show(truncate =
False)
```

► (2) Spark Jobs

```

+-----+-----+
-----+-----+
|average|title|
|count|movieId|
```

```

+-----+-----+-----+
-----+-----+-----+
|4.254283139357548|Shawshank Redemption, The (1994)
      |12513|318      |
|4.22503672482484 |Third Man, The (1949)
      |1317 |1212      |
|4.195024231735141|Godfather, The (1972)
      |8345 |858      |
|4.175523698708929|Seven Samurai (Shichinin no samurai) (1954)
      |2359 |2019      |
|4.17406071068002 |Rear Window (1954)
      |3538 |904      |
|4.173247438199274|Usual Suspects, The (1995)
      |9424 |50      |
|4.172001171513332|Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)
      |1322 |922      |
|4.170191933998413|Casablanca (1942)
      |4949 |912      |
|4.152172995341341|Dr. Strangelove or: How I Learned to Stop Worrying and Love t
he Bomb (1964)|4591 |750      |
|4.140182386042633|Schindler's List (1993)
      |9804 |527      |
+-----+-----+-----+
-----+-----+-----+

```

Command took 38.34 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 3:07:26 PM on

movies\_project  
Cmd 74

```

1 Rec_movies_gt_500_reviews = spark.sql("select Recommended_movies.movieId,
2   Recommended_movies.title, Recommended_movies.count,
3   Recommended_movies.average,predictions_with_my_ratings_table.prediction \
4   FROM Recommended_movies \
5   JOIN predictions_with_my_ratings_table ON
6   predictions_with_my_ratings_table.movieId=Recommended_movies.movieId \
7   ORDER BY predictions_with_my_ratings_table.prediction DESC\
8   LIMIT 20")
9
10 Rec_movies_gt_500_reviews.show(truncate = False)

```

### ► (3) Spark Jobs

```

+-----+-----+-----+-----+-----+
-----+-----+-----+
|movieId|title                                     |count|average
      |prediction|
+-----+-----+-----+-----+-----+
-----+-----+-----+

```

Command took 1.54 minutes -- by meghana.rwgsql@gmail.com at 8/31/2017, 3:21:29 PM on  
movies\_project  
Cmd 75

```
1 display(Rec_movies_gt_500_reviews)
```

► (3) Spark Jobs

movieId	title
1035	Sound of Music, The (1965)
1354	Breaking the Waves (1996)
913	Maltese Falcon, The (1941)
327	Tank Girl (1995)
1923	There's Something About Mary (1998)
1721	Titanic (1997)
39	Clueless (1995)
39	Clueless (1995)
2700	South Park: Bigger, Longer and Uncut (1999)



Command took 1.78 minutes -- by meghana.rwgsq@gmail.com at 8/31/2017, 3:53:03 PM on

movies\_project  
Cmd 76

```
1 #Saving the model
2 from pyspark.mllib.recommendation import MatrixFactorizationModel
3 import os
4
5 ## Mount S3 bucket nycdsabootcamp to the Databricks File System
6 s3Path = "s3a://{0}:{1}@{2}".format("AKIAI2P5MSE02JYXJVQQ",
7
8     "YJboxXSbraX4rg17aqtI+HmBjWCcpu4dxv2HW+bm",
9     "nycdsabootcamp")
9 mntPath = "/mnt/data/"
10 model_path = os.path.join('.', 'models', 'movie_lens_als')
11 # Save and load model
12 model.save(model_path)
13 same_model = MatrixFactorizationModel.load(sc, model_path)
14
```

IllegalArgumentExcepion: u'Path must be absolute: dbfs://../models/movie\_lens\_als'

Command took 0.13 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 3:50:20 PM on

movies\_project  
Cmd 77

```

1 #MODEL TUNING AND CROSS VALIDATION
2
3 from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
4
5 # 1. you first need to build a parameter grid based on your ALS model
6
7 paramGrid = ParamGridBuilder() \
8     .addGrid(als.rank, [10, 30]) \
9     .addGrid(als.maxIter, [10, 20]) \
10    .build()
11
12 # 2. use your ALS estimator and parameter grid to build cross validator
13 crossval = CrossValidator(estimator=als,
14     estimatorParamMaps=paramGrid,
15     evaluator=evaluator,
16     numFolds=3) # use 3+ folds in practice
17
18 # Run cross-validation, and choose the best set of parameters.
19 cvModel = crossval.fit(training)

```

#### ► (6) Spark Jobs

Command took 1.11 hours -- by meghana.rwgsq@gmail.com at 8/31/2017, 11:29:07 AM on

movies\_project  
Cmd 78

```

1 # Make predictions on test documents. cvModel uses the best model
2 cv_predictions = cvModel.transform(test)

```

Command took 0.07 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:38:19 PM on

movies\_project  
Cmd 79

```

1 cv_predictions.select("userId", "movieId", "rating", "prediction").show(10)
2

```

#### ► (1) Spark Jobs

```

+-----+-----+-----+-----+
|userId|movieId|rating|prediction|
+-----+-----+-----+-----+
| 92852|    148|    3.0| 2.5335486|
| 81218|    148|    1.0| 2.7732363|
| 91782|    148|    3.0| 2.9888206|
| 13170|    148|    3.0| 0.0466154|
|  1259|    148|    5.0| 3.2165732|
| 44882|    148|    4.0| 2.3504906|
| 94994|    148|    4.0| 3.119975|
| 90757|    148|    3.0| 2.9625764|

```

3673	148	2.0	2.6216304
64843	148	3.5	2.630504

only showing top 10 rows

Command took 34.71 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:38:24 PM on

movies\_project  
Cmd 80

```
1 cv_predictions = cv_predictions.dropna()
2 cv_predictions.registerTempTable("cv_predictions_table")
```

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:39:12 PM on

movies\_project  
Cmd 81

```
1 %sql select user_id, movie_id, rating, prediction from predictions
2 spark.sql("select userId, movieId, rating, prediction from
  cv_predictions_table").show()
```

#### ► (1) Spark Jobs

userId	movieId	rating	prediction
92852	148	3.0	2.5335486
81218	148	1.0	2.7732363
91782	148	3.0	2.9888206
13170	148	3.0	0.0466154
1259	148	5.0	3.2165732
44882	148	4.0	2.3504906
94994	148	4.0	3.119975
90757	148	3.0	2.9625764
3673	148	2.0	2.6216304
64843	148	3.5	2.630504
81300	148	1.0	2.3215954
118205	148	3.5	2.8588994
109121	148	4.0	3.454861
68360	148	2.0	2.942938
132268	148	2.0	1.5295749
30699	148	3.0	2.1022847
28361	148	4.0	4.177149
36723	148	1.0	1.9845809
9084	148	2.0	3.0726907
66709	148	5.0	3.728733


only showing top 20 rows

Command took 35.14 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:40:37 PM on

movies\_project  
Cmd 82

```
1 display (spark.sql("select rating, prediction from cv_predictions_table"))
```

► (1) Spark Jobs



Showing sample based on the first 1000 rows.



Command took 34.90 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:41:18 PM on

movies\_project  
Cmd 83

```
1 #Calculating root mean square
2 cv_rmse = evaluator.evaluate(cv_predictions)
3 print("Root-mean-square error = " + str(cv_rmse))
```

► (1) Spark Jobs

Root-mean-square error = 0.804667235236

Command took 38.75 seconds -- by meghana.rwgsq@gmail.com at 8/31/2017, 12:42:07 PM on

movies\_project  
Cmd 84

```
1 displayHTML("<h4>The Root-mean-square error is %s</h4>" % str(cv_rmse))
```

## The Root-mean-square error is 0.804667235236

Command took 0.07 seconds -- by meghana.rwgsql@gmail.com at 8/31/2017, 12:43:23 PM on

movies\_project  
Cmd 85

1

Cmd 86

### ► (2) Spark Jobs

movieId	title	cv_Predicted_ratings
129401	Kevin Smith: Sold...	5.5531861782073975
79507	Amar Akbar Anthon...	5.346615791320801
111329	Memorial Day (2011)	5.2839155197143555
110173	Wolf (2013)	5.074623107910156
69609	Holly (2006)	5.021418571472168
78064	Ween Live in Chic...	5.015072345733643
116155	Still Life (2013)	4.957600116729736
93729	Pageant (2008)	4.924448490142822
26433	Rockers (1978)	4.919745445251465
72866	Sandra of a Thous...	4.883637428283691
26968	Cremaster 5 (1997)	4.876115322113037
73139	Out 1: Spectre (1...	4.826381206512451
112577	Willie & Phil (1980)	4.805362701416016
79842	For Neda (2010)	4.802628517150879
98221	Year One, The (L'...	4.788150787353516
122290	Homeboy (1988)	4.785171031951904
26109	Crooks in Clover ...	4.757693290710449
123107	The Phantom of th...	4.7281880378723145
84248	Organizer, The (I...	4.706189155578613
83439	Toys in the Attic...	4.702956676483154

Command took 43.60 seconds -- by meghana.rwgsql@gmail.com at 8/22/2017, 5:46:31 PM on

Movies\_Project